Trust, Evidence, and the Limits of Automated Verification

Reflection on Homework 2

Before doing this assignment, I thought CV verification was mostly a matching problem: find the right LinkedIn profile and compare the dates, employer names, and education. The exercise made the task feel less mechanical. The difficult part was not only finding information. It was deciding what to trust, what to ignore, and how much confidence a single number should be allowed to communicate.

The document is both evidence and an input channel

The CV is supposed to be the thing the agent checks, but it can also try to influence the checker. While working on Task 2, I could put an instruction in the PDF asking a reviewer to give the document a favourable score. That makes the ordinary distinction between "instructions" and "data" much less tidy. A human reader usually notices when a CV starts talking to them. A model may read the same sentence as part of the task it has been asked to perform.

I would not treat the CV as useless or automatically false. It is a low-confidence statement of the candidate's own claims: useful for identifying what should be checked and for narrowing a search, but not sufficient by itself to settle an important fact. Profile data can then raise or lower confidence in specific claims. A CV that says it contains a skill, for example, is a reason to check that skill, not proof that it exists.

A score can look firmer than the evidence behind it

The assignment asks for one score between 0 and 1. That is convenient for grading, but it hides a lot. To produce that score, the agent has to decide which person is the candidate, whether two job titles mean the same thing, whether a missing skill matters, and whether a date difference is a real contradiction. A result just above 0.5 can be treated exactly like a result near 1.0 even though the reasoning may be much less secure.

I saw this directly with people who shared the same name, city, and industry. The first search result was not enough. One candidate could look plausible from the headline, while another matched the CV's employers, dates, degree, and skills. If this were a real KYC or recruitment process, I would want the score to come with the selected profile, the fields that matched, and the fields that did not. The score should help decide where to look next, not pretend that uncertainty has disappeared.

Optimising for the test is not the same as being truthful

Task 2 made this point more sharply. The public rules say that job descriptions are not used to determine a discrepancy. That creates a narrow place to put an embellishment without changing a

company name, job title, date, school, degree, or skill. It is a useful red-team exercise, but it is also uncomfortable: the attack works by learning what the verifier does not look at.

This is the part that reminded me of Goodhart's law. Once the reliability score becomes the target, a candidate can optimise for the score instead of for accuracy. A polished document can be more successful than an honest but messy one. A better system would make its coverage visible, record unverified claims, and be careful about treating absence of evidence as evidence of honesty.

More checking also has a cost

My first instinct was to make the verifier more cautious: search more candidates, make more tool calls, and repeat the judgement. That helped in some cases, but it also made the process slower and sometimes gave the model more opportunities to get distracted by similar profiles. The shared MCP server limit made this trade-off concrete. Reliability is not simply maximised by adding more agent steps.

The version I would trust most in practice is not one that acts as if it can resolve every case. It would compare the high-impact fields deterministically where possible, stop when one profile has a strong structured match, and send genuinely ambiguous cases to a person. That feels less impressive than full automation, but more honest about what the system knows.

The main lesson I take from this homework is that verification agents need humility. They can reduce repetitive work, but they should expose their evidence and their limits rather than hide both behind a confident-looking score.